



German University in Cairo
Faculty of Media Engineering and Technology
CSEN604: Database II
Project 1

Instructor: Dr. Wael Abouelsaadat

TAs: Belal Medhat, Ahmed Hesham, Nouran Sadek, Basem Mikhail

Release Date: April 11th, 2021

Submission: April 30th, and May 10th

Weight & Marking

The course projects are worth 20%. Project 1 (this one) is worth 14%. No best policy. The project marking scheme is included in the last page of this document.

Team Formation

This assignment should be done in teams of THREE to FIVE. You can make a team cross-tutorial and cross-major. If you cannot find a team, contact your TA **ASAP**, and you will be synched with other students.

Environment

You must use Java to develop this project. You can use any IDE of your choice such as Eclipse, IntelliJ, NetBeans, etc...

Submissions

Submissions will be made electronically via MET website.

Submission Milestones

There are two submissions in this project as listed below. You are required to make each submission on time/date. Below is a description of each submission.

Submission #	Date	Submission Focus	Methods to be implemented (or changed) in DBApp.java
1	April 30 th , 2021, 11:50PM	create, insert update, delete	createTable insertIntoTable updateTable deleteFromTable
2	May 10 th , 2021, 11:50PM	Grid Index + Grid Index integration	createIndex insertIntoTable selectFromTable deleteFromTable updateTable

Note that you will resubmit methods such as **updateTable** twice. The first submission without an index implementation. In the second submission, you will modify the code to include handling indices.

Description

In this project, you are going to build a small database engine with support for Grid Index. The required functionalities are 1) creating tables, 2) inserting tuples, 3) deleting tuples, 4) searching in tables linearly, 5) creating a Grid index, and 6) searching using the Grid index.

Note that this is a simplified data base engine – for example, you are neither required to support foreign keys nor referential integrity constraints. Only implement what is mentioned in this doc.

The description below is numbered for ease of communication between you and course staff.

Tables

- 1) Each table/relation will be stored as pages on disk.
- 2) Supported type for a table's column is one of: [java.lang.Integer](#), [java.lang.String](#), [java.lang.Double](#), [java.util.Date](#) (Note: date acceptable format is "YYYY-MM-DD")

Pages

- 4) A page has a predetermined fixed maximum number of rows (N). For example, if a table has 40000 tuples, and if $N=200$, the table will be stored in 200 binary files.
- 5) You are required to use Java's binary object file (.class) for emulating a page (to avoid having you work with file system pages, which is not the scope of this course). A single page must be stored as a serialized Vector ([java.util.Vector](#) because Vectors are thread safe). Note that you can save/load any Java object to/from disk by implementing the [java.io.Serializable](#) interface. You don't actually need to add any code to your class to save it the hard disk. For more info, check this: https://www.tutorialspoint.com/java/java_serialization.htm
- 6) A single tuple should be stored in a separate object inside the binary file.
- 7) You need to postpone the loading of a page until the tuples in that page are actually needed. Note that the purpose of using pages is to avoid loading the entire table's content into memory. Hence, it defeats the purpose to load all pages upon program startup.
- 8) If all the rows in a page are deleted, then you are required to delete that page. Do not keep around completely empty pages. In the case of insert, if you are trying to insert in a full page, shift one row down to the following page. Do not create a new page unless you are in the last page of the table and that last one was full.
- 9) You might find it useful to create a Table java class to store relevant information about the pages and serialize it just like you serialize a page. Note that to prevent serializing an attribute, you will need to use the Java [transient](#) keyword in your attribute declaration. Read more here: <https://www.tutorialspoint.com/difference-between-volatile-and-transient-in-java>

Meta-Data File

- 10) Each user table has meta data associated with it; number of columns, data type of columns, which columns have indices built on them.

11) You will need to store the meta-data in a text file. This structure should have the following layout:

Table Name, Column Name, Column Type, ClusteringKey, Indexed, min, max

ClusteringKey is set true if the column is the primary key. For simplicity, you will always sort the rows in the table according to the primary key. That's why, it is called the clusteringkey. Only 1 clustering key per table.

min and max refer to the minimum and maximum values possible for that column.

For example, if a user creates a table/relation CityShop, specifying several attributes with their types, etc... the file will be:

```
Table Name, Column Name, Column Type, ClusteringKey, Indexed, min, max
CityShop, ID, java.lang.Integer, True, False,0,10000
CityShop, Name, java.lang.String, False, False, "A", "ZZZZZZZZZZ"
CityShop, X, java.lang.Double, False, True, 0,1000000
CityShop, Y, java.util.Double, False, True, 0,1000000
CityShop, Z, java.lang.Double, False, True, 0,1000000
CityShop, Specialization, java.lang.String, False, True, "A", "ZZZZZZZZZZ"
CityShop, Address, java.lang.String, False, false, "A", "ZZZZZZZZZZ"
```

The above meta data teaches that there are 1 table of 7 tuples (ID, Name, X,Y,Z, Specialization, Address). There are 4 indices created on this table CityShop. Note that if a multidimension index has been created, you will not be able to know which columns are used with which from this file. You will only know if there is an index.

12) You must store the above metadata in a single file called **metadata.csv**. Do not worry about its size in your solution.

13) You must use the metadata.csv file to learn about the types of the data being passed and verify it is of the correct type. So, do not treat metadata.csv as decoration! ☺

14) You can (but not required to) use reflection to load the data type and also value of a column, for example:

```
strColType = "java.lang.Integer";
strColValue = "100";
Class class = Class.forName( strColType );
```

```

Constructor constructor = class.getConstructor( ... );
... = constructor.newInstance( );


```

For more info on reflection, check this article:

<http://download.oracle.com/javase/tutorial/reflect/>

Indices

15) You are required to use Grid Index to support creating primary and secondary indices. A grid index can be single dimension (in this case a 1D array), or multidimensional one (multidimensional array). Your solution should support any number of dimensions. The Grid index should be implemented as below illustration from lecture 5 slides. Each arrow coming out from a bucket is the address on disk of the table's page which contains the row of interest. Similarly, each arrow coming from a cell in the in-memory grid is the address on disk of an index's bucket/page.



16) Each column has an associated range (minimum, maximum). You should use this range to create the divisions on the Grid Index assigned scale. You should target create close to 10 divisions if possible. For example, in the case of a String with min: "a", max: "z", the divisions can be: a-c, d-f, g-i, j-l, m-o, p-q, r-s, t-u, v-w, x-z

17) Once a table is created, you do not need to create any page until the insertion of the first row/tuple.

18) You should update existing relevant indices when a tuple is inserted/deleted.

19) If a secondary index is created after a table has been populated, you have no option but to scan the whole table.


```

// following method creates one table only
// strClusteringKeyColumn is the name of the column that will be the primary
// key and the clustering column as well. The data type of that column will
// be passed in htblColNameType
// htblColNameValue will have the column name as key and the data
// type as value
// htblColNameMin and htblColNameMax for passing minimum and maximum values
// for data in the column. Key is the name of the column
public void createTable(String strTableName,
                      String strClusteringKeyColumn,
                      Hashtable<String,String> htblColNameType,
                      Hashtable<String,String> htblColNameMin,
                      Hashtable<String,String> htblColNameMax )
                                throws DBAppException

// following method creates one index - either multidimensional
// or single dimension depending on the count of column names passed.
public void createIndex(String strTableName,
                      String[] strarrColName) throws DBAppException

// following method inserts one row only.
// htblColNameValue must include a value for the primary key
public void insertIntoTable(String strTableName,
                          Hashtable<String,Object> htblColNameValue)
                                throws DBAppException

// following method updates one row only
// htblColNameValue holds the key and new value
// htblColNameValue will not include clustering key as column name
// strClusteringKeyValue is the value to look for to find the rows to update.
public void updateTable(String strTableName,
                      String strClusteringKeyValue,
                      Hashtable<String,Object> htblColNameValue )
                                throws DBAppException

// following method could be used to delete one or more rows.
// htblColNameValue holds the key and value. This will be used in search
// to identify which rows/tuples to delete.
// htblColNameValue enteries are ANDED together
public void deleteFromTable(String strTableName,
                          Hashtable<String,Object> htblColNameValue)
                                throws DBAppException

public Iterator selectFromTable(SQLTerm[] arrSQLTerms,
                              String[] strarrOperators)
                                throws DBAppException

```

Here is an example code that creates a table, creates an index, does few inserts, and a select;

```

String strTableName = "Student";

DBApp dbApp = new DBApp( );

Hashtable htblColNameType = new Hashtable( );
htblColNameType.put("id", "java.lang.Integer");

```

```

htblColNameType.put("name", "java.lang.String");
htblColNameType.put("gpa", "java.lang.double");
dbApp.createTable( strTableName, "id", htblColNameType /* in addition to min
max hashtables" );
dbApp.createIndex( strTableName, new String[] {"gpa"} );

Hashtable htblColNameValue = new Hashtable( );
htblColNameValue.put("id", new Integer( 2343432 ));
htblColNameValue.put("name", new String("Ahmed Noor" ) );
htblColNameValue.put("gpa", new Double( 0.95 ) );
dbApp.insertIntoTable( strTableName , htblColNameValue );

htblColNameValue.clear( );
htblColNameValue.put("id", new Integer( 453455 ));
htblColNameValue.put("name", new String("Ahmed Noor" ) );
htblColNameValue.put("gpa", new Double( 0.95 ) );
dbApp.insertIntoTable( strTableName , htblColNameValue );

htblColNameValue.clear( );
htblColNameValue.put("id", new Integer( 5674567 ));
htblColNameValue.put("name", new String("Dalia Noor" ) );
htblColNameValue.put("gpa", new Double( 1.25 ) );
dbApp.insertIntoTable( strTableName , htblColNameValue );

htblColNameValue.clear( );
htblColNameValue.put("id", new Integer( 23498 ));
htblColNameValue.put("name", new String("John Noor" ) );
htblColNameValue.put("gpa", new Double( 1.5 ) );
dbApp.insertIntoTable( strTableName , htblColNameValue );

htblColNameValue.clear( );
htblColNameValue.put("id", new Integer( 78452 ));
htblColNameValue.put("name", new String("Zaky Noor" ) );
htblColNameValue.put("gpa", new Double( 0.88 ) );
dbApp.insertIntoTable( strTableName , htblColNameValue );

SQLTerm[] arrSQLTerms;
arrSQLTerms = new SQLTerm[2];
arrSQLTerms[0].__strTableName = "Student";
arrSQLTerms[0].__strColumnName= "name";
arrSQLTerms[0].__strOperator  = "=";
arrSQLTerms[0].__objValue    = "John Noor";

arrSQLTerms[1].__strTableName = "Student";
arrSQLTerms[1].__strColumnName= "gpa";
arrSQLTerms[1].__strOperator  = "=";
arrSQLTerms[1].__objValue    = new Double( 1.5 );

String[]strarrOperators = new String[1];
strarrOperators[0] = "OR";
// select * from Student where name = "John Noor" or gpa = 1.5;
Iterator resultSet = dbApp.selectFromTable(arrSQLTerms , strarrOperators);

```


- 25) For the parameters, the name documents what is being passed – for example `htblColNameType` is a hashtable with *key* as `ColName` and *value* is the `Type`.
- 26) Operator Inside `SQLTerm` can either be `>`, `>=`, `<`, `<=`, `!=` or `=`
- 27) Operator between `SQLTerm` (as in `strarrOperators` above) are **AND, OR, or XOR**.
- 28) **`DBAppException`** is a generic exception to avoid breaking the test cases when they run. You can customize the Exception by passing a different message upon creation. You should throw the exception whenever you are passed data you that will violate the integrity of your schema.
- 29) **`SQLTerm`** is a class with 4 attributes: `String _strTableName`, `String _strColumnName`, `String _strOperator` and `Object _objValue`
- 30) Iterator is java.util.Iterator It is an interface that enables client code to iterate over the results row by row. Whatever object you return holding the result set, it should implement the Iterator interface.
- 31) You should check on the passed types and do not just accept any type – otherwise, your code will crash will invalid input.
- 32) You are not supporting SQL Joins in this mini-project.

Directory Structure

- 33) Your submission should follow the given maven project template (available on cms/MET)
- 34) In the resources directory, include a configuration file *DBApp.config* which holds a parameters as key=value pairs

```
MaximumRowCountinTablePage    = 200
MaximumKeysCountinIndexBucket = 100
```

Where

`MaximumRowCountinTablePage` as the name
indicates specifies the maximum number
of rows in a page.

While `MaximumKeysCountinIndexBucket` is the
max keys you can store in a grid index
bucket.

Note: you can add any other configuration parameter in
`DBApp.config`

35) *DBApp.config* file could be read using [java.util.Properties](https://docs.oracle.com/javase/7/docs/api/java/util/Properties.html) class

36) If you want to add any external library, you should add it to the pom.xml file in maven as a dependency. Do not include any external 3rd jar files in your submission.

Bonus – Worth 1% of course grade

37) Add support for processing SQL statements. For that, you will need a SQL parser. You can use ANTRL, a parser generator. Here is a page which include a number of ready to use SQL grammars: <https://wwwantlr3.org/grammar/list.html>

38) If you do the bonus, only accept statements that you can run in your mini database engine. For example: create constraint Should throw an exception.

39) To pass a SQL string, such as create table ..., add a method with the following syntax:

```
// below method returns Iterator with result set if passed
// strbufSQL is a select, otherwise returns null.
public Iterator parseSQL( StringBuffer strbufSQL ) throws
DBAppException
```

Marking Scheme

1) +5 Each table/relation will be stored as binary pages on disk and not in a single file

2) +2 Table is sorted on key

3) +4 Each table and page is loaded only upon need and not always kept in memory.

Page should be loaded into memory and removed from memory once not needed.

4) +2 A page is stored as a vector of objects

5) +3 Meta data file is used to learn about types of columns in a table with every
select/insert/delete/update

6) +2 Page maximum row count is loaded from metadata file (N value)

7) +3 A column can have any of the 6 types

8) +1 TouchDate supported

9) +6 select without having any index created is working fine

10) +8 select with the existence of an index that could be used to reduce search space

- 11) +6 insert without having any index created is working fine
- 12) +8 insert with the existence of an index that could be used to reduce search space
- 13) +6 delete without having any index created is working fine
- 14) +8 delete with the existence of an index that could be used to reduce search space
- 15) +6 update without having any index created is working fine
- 16) +8 update with the existence of an index that could be used to reduce search space
- 17) +6 creating an index correctly for a given column whether key column or otherwise.
- 18) +4 saving and loading index from disk
- 19) +12 Inserting and deleting from index correctly.

Total: /100

Other Deductions:

- Not respecting specified method signatures → -5
- Not respecting specified directory structure → -2
- Not submitting ant, maven or make file → -1
- not performing binary search on table pages → -3
- Using ArrayList instead of vector → -2